

CS 134 – Distributed Systems

Spring 2025 – Professor Konstantinos Kallas

Lecture 1: What is a Distributed System?

Distributed System: Multiple interconnected computers that cooperate to provide some service

Challenges in Distributed Design

- Partial Failures
- Ambiguous Failures (we don't know what the errors are)
 - Server failure, Network failure, or both
- Concurrency
 - Ensure consistency of distributed state

Lecture 2: MapReduce, Clock Synchronization

MapReduce (and Hadoop)

- Goals
 - Scalability: Throughput grows with number of machines
 - Fault Tolerance: Makes progress despite machine failures
 - No latency requirement: MapReduce is used for batch jobs
- Naïve Approach
 - Idea: Partition the data, split across servers. Each server is responsible for own data
 - Issue: We may need to access data split across machines
- MapReduce Algorithm
 - Partition: Split dataset into partitions
 - Map: Outputs Key Value Pairs
 - Coalesce: Organize Key Value Pairs into new partitions
 - Reduce: Aggregation using a reducer function
- Issues
 - Synchronization barrier: Cannot run reduce tasks until all map tasks finish

Paradigm: Fault Tolerance via Master Node

- To achieve fault tolerance, the master node:
 - Assigns map/reduce jobs to workers
 - Keeps track of status of tasks
 - Keeps track of machine failures (by retransmitting)
- Abstractions
 - Control (fork) vs Data Communications
 - Client-Server
 - Client makes requests to servers
 - Server listens for requests, processes and returns responses
 - Issue: Cannot Scale
 - Client-Server with threading

- Listener thread spawns thread to handle each request it receives
 - Listener thread puts requests in a queue. Thread pool is fixed
- RPC (Remote Procedure Calls)
 - Server offers API of methods with signatures
 - Clients invoke a method using args, waits for responses
- Issue: If master fails, we lose all state of task execution status

Checkpointing

- Issue: If any partition fails, reduce stage will never start
- Solution: Checkpoint after each phase to fall back
 - Slow
 - Writes to disk → Expensive
 - Orders of magnitude worse than writing to memory
- Apache Spark
 - Operations completed in RAM, changes only persist after commits

Replicated State Machines

- Idea: Replicate master state to tolerate failures
- Issue: Inconsistent replicas
 - Workers can potentially communicate faster than the masters communicate to each other:
 - W2 finishes task 2 before M1 tells M2 that W1 has been assigned task 3
- Solution: Synchronization
 - Idea: Apply updates in the same order across all replicas
 - Approach: Given list of received updates, order them based on time of receipt
 - Issue: Clocks are not in sync across replicas
 - Store the sequence of transitions to construct the state
 - Abstractly defined as a state machine
 - Approaches:
 - Replication over Time:
 - Start new replica when one goes down

- Assumes execution state from failed machine is available to continue from
- Replication over Space:
 - Run multiple replicas simultaneously
 - Assume that when something fails, all state is lost

Clock Synchronization:

- Naïve Approaches
 - GPS broadcasting (accurate ~1ms) → Power hungry, doesn't work indoors
 - Correct skew based on reference clock → Clock drift, unbounded network delays
 - Potential workaround using Network Time protocol
- Cristian's Algorithm
 - Assume time server is available
 - Client sends request packet using local timestamp
 - Server timestamps its receipt of client's request using its own local time
 - Server responds with packet containing its local time on receive and send
 - Client locally timestamps receipt of server's response of send
 - Round Trip Time = $\hat{\partial}_{\text{req}} + \hat{\partial}_{\text{resp}} = (T_4 - T_1) - (T_3 - T_2)$

Lecture 3: Clock Properties, Lamport Clocks, Vector Clocks

Desired Ordering Properties

- Idea: We just want the order between events to be the same
 - All replicas agree on an order:
 - For all i, j , $C_i(a) < C_i(b)$ iff $C_j(a) < C_j(b)$
 - Total Order: for all a, b , $C(a) < C(b)$ or $C(b) < C(a)$
 - Accuracy: If a happened before b , $C(a) < C(b)$
- $C_i(a) < C_i(b) \rightarrow a$ is ordered before b

Causality

- If A caused B , then A happened before B
- Synchronous Model: All messages are delivered in less time than N
- Asynchronous: A message can take an arbitrary amount of time to be delivered

Logical Clocks

- Idea: Associate every event with a logical time
 - Across Processes, Transitivity, In same process
- Allows all replicas in an RSM to execute updates in the same order
 - Note: cannot handle failures (using only logical clocks)

Lamport Clocks

- Idea:
 - Associate all events with a clock time, order based on clock time
- Algorithm
 - Each process p_i maintains a local clock C_i
 - All process clocks start at time 0
- Issue: We cannot tell if an event occurs before another
 - 1, 2 can occur before 1,2,3,4,5
- Issue: Concurrency
 - If events A, B are concurrent, we cannot establish a total order
 - Solution: Break ties in the same manner across all replicas

- i.e. Append the process number to each event
- Process P_i timestamps event e with $C_i(e) \cdot i$
 - $C(a) \cdot i < C(b) \cdot j$ when:
 - $C(a) < C(b)$ OR $C(a) == C(b)$ AND $i < j$

RSMs with Lamport Clocks

- Idea: Place events into a local queue sorted by increasing $C(x)$
 - We only execute a process if we receive an ACK from all other processes that is greater (given causality)
- Optimizations
 - Reducing waiting
 - Periodically ping all other nodes to sync clock
 - Tolerate temporary inconsistency
 - Apply update immediately upon receiving it
 - Maintain log of updates to roll back
 - Use causal ordering rather than total ordering
 - Not all events require ordering with everyone else
 - Issue: If any one replica goes down, all other replicas are unable to make progress

Vector Clocks

- Idea: Determine a precise partial order of events in distributed system
- Algorithm
 - Label each event with a vector $[c_1, c_2, \dots, c_i]$
 - C_i is a count of events in process I that causally precede e
 - All vectors start at $[0, 0, \dots, 0]$
 - Updating a vector:
 - For local event on same process, increment c_i by 1
 - If receiving a message from another process with vector $[d_1, d_2, \dots, d_i]$:
 - Modify local vector to $c_k = \max\{c_k, d_k\}$
 - Increment local entry c_j in vector by 1
- Comparing Vector clocks

- $V(a) == V(b) \rightarrow a_k == b_k$ for all k
- $V(a) < V(b) \rightarrow a_k \leq b_k$ for all k and $V(a) \neq V(b)$
- Concurrency check:
 - $A \parallel b \rightarrow a_i < b_i$ and $a_j > b_j$ for some i, j
- $V(a) < V(z)$ if and only if there is a chain of events linked by \rightarrow between a and z

Lecture 4: Failures, Chandy-Lamport Algorithm

Logical clocks break on failures

Fault Tolerance

- A system should be consistent and should guarantee that everyone will eventually see updates
 - A system S satisfies its specification $Spec$ given that some classes of faults happen to it

Failures

- Crash failures
 - Component is unreachable and never comes back to life
 - We cannot reassign work back to original thread, another process takes over
- Network failures
 - Messages are dropped or reordered
 - Solved through retransmission and attaching identifiers to deduplicate
- Byzantine failures
 - Modified or corrupted messages or processes
 - Component does not follow protocol and executes arbitrarily

Recovering from failures: Checkpoint Storage

- Algorithm
 - Every N seconds
 - Store checkpoints on persistent storage (disk)
 - If a component crashes:
 - Retrieve the latest checkpoint
 - Restore the component state based on it
 - Continue execution from that point on
- Issues
 - Domino Effect: Uncoordinated checkpointing across processes
 - Solution: Coordinated Checkpointing/Chandy-Lamport

- Inconsistency if a message is transferred between checkpoint executions

Global Checkpointing (Uncoordinated)

- Idea: Take a coordinated snapshot of a distributed system asynchronously
 - Useful for garbage collection, deadlock detection, and evaluating stable properties
- Token Ring
 - Idea: Rotate lock possession across processes in circle/ring order
 - Issue: Snapshots may result in multiple processes thinking they have the lock after recovery
 - P_1 and P_2 both have receives but no sent
 - Note: In weak guarantees that use one lock, there is no consistency issues

Global Snapshots (Uncoordinated)

- Idea: captured state reflects causality in regards to state of individual processes and every communication channel (things sent that are not received)
- Issue: Some snapshots may make the locks unrecoverable
 - P_2 thinks lock is sent, P_3 has not received

Strawman (Coordinated Screenshots)

- Idea: When a snapshot message is received, all processes stop working and prepares for a snapshot
 - High cost on execution
 - Consistent

FIFO Channels

- Idea: No message sent before another arrives AFTER another (using TCP)

Chandy-Lamport Snapshots

- Idea: We need to handle crash failures
- Implementation
 - N processes in system
 - Assume processes cannot fail during snapshot

- Any process can initiate a snapshot
 - Two FIFO unidirectional channels between each process pair
 - Assume no message duplication or loss
- Core Ideas
 - After taking a checkpoint, send a broadcasting marker message to other processes to also take checkpoints
 - Marker enables processes to discover the need for taking a snapshot
 - Marker servers as a “barrier” on every network channel that all messages received before the marker are included
- Algorithm
 - Note that in this class, we capture using the SENDS of a message
 - Let process p_k initiate the snapshot
 - p_k records its local state (takes a checkpoint)
 - p_k sends out marker messages on every outgoing network channel
 - p_k starts recording messages on every incoming channel
 - Let p_i receive the marker message from p_k (via channel c_{ki})
 - If marker is first marker p_i has seen
 - p_i records its state
 - p_i marks channel c_{ki} as empty
 - p_i sends a marker out to every outgoing channel c_{ij}
 - p_i starts recording incoming messages on all channels other than c_{ki}
 - If p_i has already seen a marker
 - p_i stops recording on channel c_{ki}
 - p_i sets c_{ki} ’s final state as the sequence between the start and end of recording
 - The snapshot is complete when every process has received marker messages from every incoming channel
 - i.e. $(\text{Processes} \cdot (\text{Processes} - 1))$ markers are sent
- Rollback
 - Recognize there is a failure in the system
 - All active nodes roll back to latest completed snapshot

- A new/restarted node recovered from its snapshot
 - Loads everything into memory
- Continue
- Issues
 - If message themselves are nondeterministic, we might lose message on rollback
 - Ex: coin flip

Nondeterminism:

- Sources: Randomness, user input, network delays, process scheduling, OS interrupts
- Issues
 - Reasoning within distributed systems (state explosion)
 - Messes with replication and recovery
- Solution: Message logging
 - Idea: between checkpoints, log all events that lead to nondeterminism
 - Log enough to ensure that same messages are sent as in original execution
 - Thus, we can now guarantee. That we can roll back and execute deterministically

Lecture 5: Fault Tolerance, Primary Backup Replication, Request Types

Fault Tolerance Trade Offs

- Cost on normal execution (latency/CPU/memory)
- Recovery freshness: how far back of a snapshot do we recover from
 - High freshness from storing state
- Time to recover
- For reference, Chandy-Lamport is non-blocking on normal execution, kind of fresh (some messages are lost), and time to recover is just ok

API Request Types and Workloads

- Service APIs: Supports many API methods, used to interace between client and server
- API Variability
 - Modifying storage and state → Reads/Writes
 - Accesses a lot of states → IO/CPU heavy
 - Accessing many items from state → implies coordination
- Workload Variability
 - Uniform load across methods
 - Read heavy vs Write heavy
 - Skewed access of a few keys
 - Load spikes/bursty traffic
- API Optimization
 - Caching works for read heavy applications (but sucks for write heavy)
 - Key value stores are assumed to be read heavy for this course
 - Batch requests if IO/network costs are high
 - Send N requests to replica at once
 - Issue: If requests arrive infrequently, we might have to wait for a long time
 - Solution: Batch timeout
 - If T seconds have passed since last batch, send an incomplete batch

- Throughput increases but latency increases too
- Pipeline to maximize CPU usage during I/O
 - Before processing requests, send requests to backup in parallel

Primary Backup Replication (Fault Tolerance)

- Idea: Two servers (primary and backup) run concurrently
 - Primary interacts with client
 - Backup copies primary's state
 - Backup replaces primary if it dies
- Big Ideas
 - If primary fails, replace with backup
 - Backups need to be aware of who the new primary is
 - Naïve Approach: Hardcode primary into the app code
 - We have the same issue
 - Approach: Use a View Service
 - If backup fails, add another machine as backup
 - Sync primary to backup before anything is externally visible/observable to the user
 - This allows for rollbacks to be possible
 - It is ok for the primary to be out of sync if the change is not externally visible
 - Resend messages that are dropped. Use random number as hash to deduplicate
 - After some time, if no response, assume backup is dead and bring in a new backup
 - We must wait for the backup to acknowledge it has received the message
 - Transferring from Primary → Backup
 - Bootstrap
 - Snapshot of primary's state (send all state)
 - Requires serialization, sending over network, then deserialization into memory

- i.e. it is slow
- Steady state
 - Send every operation and all sources of nondeterminism
 - Includes external inputs/requests
- Issues
 - Sending a response and it fails
 - Client deduplication
 - Client has data which backup does not have
- Optimizations
 - Primary does not need to consult backup for all operations (specifically reads)
 - This is only OK if the backup is externally consistent with the primary

Implementing Primary Replicas

- Naïve approach: RSMs with logical clocks
 - Idea: Implementation is oblivious to clocks and uses a library to keep things in sync
 - Functions receive messages from client and syncs with backups before sending a response
 - Issue: Does not capture nondeterminism in execution
- Virtual Machines
 - Idea: Isolate application, monitor nondeterminism
 - Primary and backup execute on two VMs (ideally on different machines)
 - Primary logs inputs and nondeterminism (interrupts)
 - Logs delete after some time
 - Backup applies inputs from log
 - Primary waits for backup to return
 - Primary and backup monitor each other
- Log-based VM replication (Virtual Machine Monitors)
 - Idea: Primary and Backup VMM
 - Sources of nondeterminism are decided upfront by VMM
 - Primary and backup never diverge

- Approach:
 - Primary VMM sends log entries to backup VMM over the logging channel
 - Backup VMM (hypervisor) replays log entries
 - Stops backup VM at next input event or nondeterministic instruction
 - Delivers same input event as primary used
 - Delivers same nondeterministic value as primary

Lecture 6: View Service and Split Brain

View Service

- Idea: Maintain current membership of primary-backup service via a view
- Big Ideas:
 - View service changes view when:
 - Primary OR any backup fails
 - Periodically send heartbeats
 - Transitioning between views
 - View Service broadcasts view change to all replicas
 - If view service dies, clients cannot find primary
 - Primary syncs with new backup if there is one
 - New primary is up to date as we only promote a previous backup
 - New primary must ACK new view once backup is up to date
 - Can become stuck if primary fails during view change
- Scalability
 - Client does not need to contact view service before every operation
 - **Cache** view across operations
 - If client gets no response/bad response from replica it thinks is primary, it invalidates its cache
- Fault Tolerance
 - View Service is not on the critical load path
 - Assuming a reliable network, we can still end up in split brain

Split brain

- Two partitions of nodes disagree on the state of the system
 - Leads to arbitrary inconsistency of the system
- Solution: Primary must forward all operations to backups
 - i.e. get ACKs from backups that they too recognize the primary
 - Backups know who is primary as only a backup can be promoted as a primary

Lecture 7: Linearizability and Consistency Models

Desired Ordering Properties

- Total Ordering of Writes
 - Abstracts application into single thread process for the user
- Reads relative to writes
 - Once a read returns a particular value, all later reads return the same value or the value of a later write
 - i.e. No taking back values after confirming
 - Once a write completes, all later reads should return the value of the write or the value of a later write
 - Ensures freshest results

Trace: Sequence of requests and responses

- A trace is sequential if no requests overlap
- A trace is linearization of another trace if all overlapping events are made non overlapping and the order of non overlapping is preserved

Linearizability (Operational)

- Properties
 - Total ordering of writes
 - Read returns the last completed write
- Single copy semantics
 - Externally visible effects of writes and reads are equivalent to if there existed a single copy
- Linearizability is always defined based on a specification
 - Given a sequential specification S , an implementation I is linearizable iff for all traces t in I , there exists a linearization t' that belongs in S

Consistency Models

- Consistency models are properties that all executions of a given system satisfy
 - Performance and Consistency are on an inverse scale

- Eventual (Easiest to write)
- Read-after-write
- Causal
 - Idea: Order of causally related writes must be preserved in values returned to reads
 - Allows for lazy synchronization between replicas
 - If $W_1 \rightarrow W_2$, if a read sees the effect of W_2 , it must see effect of W_1
- Sequential
 - Idea: There exists a sequential order of all events that can explain the execution
 - Read time order does not matter for independent clients
 - Same as linearizability but allows arbitrary reordering across clients
 - Global sequential order of all events across all clients
 - Requests within a client cannot be reordered
 - i.e. a global order of writes
 - Implemented using Lamport Clocks or having single primary doing all writes
 - Increases blocking and performance overhead
 - Does not hold when clients are
 - No longer independent (use shared resources like channels)
 - creating diverging states
 - Should be prevented by the global order
- Linearizability (Most consistent, more expensive)

Lecture 8: Consensus, Safety, Liveness, Paxos

Safety and Liveness

- Systems can be defined via a set of execution traces
- Systems satisfy a property P if all of its execution traces satisfy a property
- We can prove a system does not satisfy a property via a counterexample trace
- Safety: System never reaches an undesirable state
 - Counterexample: a finite trace
 - Easier to prove, check, and guarantee
- Liveness: System makes progress eventually
 - Counterexample: an infinite trace

Consensus

- Issue: RSM is not tolerant to network partitions. Operations block when some machines are unavailable
- Idea: Apply update only if majority of replicas commit to it
 - If $2f + 1$ replicas, we need $f + 1$ to commit
 - Higher tolerance if there are more replicas
 - $2f + 1 \rightarrow$ Avoids ambiguous acceptance of other updates
- Assume replicas start in sync with one another
- Properties in Consensus
 - Termination: Each correct process eventually decides on a value (Liveness)
 - Race Condition prevents liveness: Between one proposer's Prepare and Accept phases, n_{proposer} is updated by another proposer
 - Solution: Proposer sends propose messages to all accepters
 - Retry with higher proposal number if majority do not accept
 - Agreement: All correct processes decide on the same value
 - Validity: The agreed upon value must be one of the proposed values
- Consensus Phases: Propose \rightarrow Decide \rightarrow Accept

Paxos

- Processes must know the number of acceptors before the start
 - Processes can persist the value they accept
 - Impossible to have all numbers accept different values due to consensus
 - For f faults, $2f + 1$ acceptors are needed
 - For f faults, $f + 1$ proposers are needed
- Pros
 - Only needs majority of replicas to be up
 - No centralized view service
- Cons
 - Needs two rounds of inter-replica communication
 - Arbitrary wait until next operation is accepted
- A process can play multiple roles
 - Proposers propose values
 - Acceptors contribute to choosing among proposed values
 - Learners learn the agreed upon value
- Idea: Protocol runs over multiple rounds
 - Once a value is accepted by a majority, a different value cannot be accepted in a later round
 - Prepare Phase
 - Proposer sends a unique proposal number to all acceptors
 - Waits for a majority commitment from acceptors
 - Accept Phase
 - Proposer sends proposed value to all acceptors
 - Waits for a majority acceptance
 - Learn Phase
 - Learners discover value accepted by the majority

Paxos (Technical version)

- Prepare
 - Create n , a proposal number (not a proposal value)
 - Proposer sends $\text{prepare}(n)$ to all acceptors

- n is unique, and higher than previous n that this proposer has used before
 - localized n allows us to avoid finding consensus over higher global n value
 - Acceptors (when receiving prepare(n))
 - Check if acceptor previously promised a higher proposal ($n_{\text{acceptor}} \leq n_{\text{proposer}}$)
 - If yes, respond with prepare-failed()
 - If no, check if acceptor has accepted any previous proposal
 - If yes, respond with promise(n , (n_{acceptor} , $\text{val}_{\text{acceptor}}$)) for previously accepted value
 - Otherwise,
 - Reply promise(n)
 - Set $n_{\text{acceptor}} = n_{\text{proposer}}$
 - The above is called the promise phase
- Accept
 - Proposer
 - When majority agree on promise(n),
 - Send accept(n , val) to acceptors
 - Acceptors (on receiving accept(n))
 - Check if acceptor previously promised a higher proposal ($n_{\text{acceptor}} \leq n_{\text{proposer}}$)
 - If yes, respond with accept-failed()
 - If no,
 - Reply accepted(n , val)
 - Broadcast accepted(n , val) to learners
 - Set $n_{\text{acceptor}} = n_{\text{proposer}}$
 - Set $\text{val}_{\text{acceptor}} = \text{val}_{\text{proposer}}$

Lecture 9: FLP, RSM with Paxos, Paxos Efficiency

Two General Problem

- Two generals on opposite sides of an invading army
 - If they both attack simultaneously, they can win
 - If they attack at different times, they all die
- They have no time to coordinate at all
 - Can only send messengers through the enemy camp
 - Messengers can be captured and killed

FLP Result (Fischer, Lynch, Paterson)

- The following situation possibly never terminates
 - Binary consensus (agree on 0 or 1)
 - Asynchronous network
 - One unreliable process (failures)

Liveness Solutions for Paxos

- When a proposal fails, back off for a random period of time before retrying
- Use a pre-determined ordering of proposers
- Elect a leader to propose requests

RSM with Paxos

- Idea: Apply all updates in a consistent order at all replicas
 - Available as long as majority is able to communicate
- Approach
 - Set of replicas (must be an odd number)
 - Replicas play all roles and agree on sequence of requests to be processed
 - Shared log of requests
 - Each replica has its own copy of the log
 - Paxos comes to consensus about each slot of the log
 - Results in performance issues
 - Requires implementation of garbage collection of log

- If another request arrives while another is pending, store it in a local queue of requests to apply

Improving RSM Paxos Performance

- Idea: Once a proposed value is accepted, jump directly into acceptance phase
 - Skips a round of communication per proposal
 - General industry practice too
- Desirability Property: Improve the best case without making the worst case worse

Leader-Based Paxos (Optimization)

- Use Paxos to elect a leader rather than selecting values via a shared sequence of views
 - Operate as primary-backup after election (all proposals sent to leader)
- Logistics
 - Leader Crash → Elect a new leader by deciding on a new view with Paxos
 - When to View Change → Make views have a validity period and make replicas send heartbeats to each other
 - Drawbacks → Leader may be far from the client

Lecture 10: Chubby, Zookeeper

Chubby

- General Info:
 - File system service for coordination
 - Each file is associated with a lock
 - Used to abstract away Paxos
 - Clients assume service doesn't fail
 - Chunk-master (leader)
 - If lock acquired, declare self as chunk-master
 - If not acquired, find the chunk-master
 - Chubby as a lock service (Service Architecture)
 - Libraries are maintained by users managing the machines
 - Services handle RPC calls and decouples deployment and operation
 - Maintained by organizations
 - Replicated Service using Paxos to implement fault-tolerant log
 - Chubby cells can have many locks
- Chubby Goals
 - High availability
 - Thousands of clients interact on small number of files
 - Event notification to reduce polling
 - Scalability (service cannot bottleneck)
 - Access control (avoids financial and security issues)
- Workflow
 - Chubby clients ping the Chubby service to acquire lock
 - If Chubby session with the lock dies, lock is released
 - Locks have timed lease to guarantee validity
 - Client sends heartbeat messages to service before expiration
 - If lease expires, lock is released
 - Assumes clock drift is bounded by $\lll 1s$
 - i.e. we assume the lock is held longer than it actually is

- Issue: Latency of storage service extends past lease expiration
 - i.e. effects (from non-Chubby services) can happen after a lock is released
- Solution: Client passes a lockID object to storage
 - Storage checks with lock service whether lock is still valid
 - i.e. Lock ID piggybacking
- Scalability
 - Caching client data read
 - Issue: Local reads from cache can return stale results (violating linearizability)
 - Solution: Master invalidates cached copies upon update
 - Master stores all clients who may have cached a data item
 - Traffic is read heavy and Chubby has no latency requirements on writes
 - Proxies
 - Proxies batch heartbeats to forward less over network

Lock Granularity

- Fine-grained locks (grab lock per action/transaction)
 - Requires higher scalability of lock service
 - Brief lock service unavailability stalls everyone
 - Clients must be prepared to lose locks frequently
- Coarse-grained locks (grab lock for minutes) (Chubby)
 - Scaling responsibility falls to clients
 - Unavailability (seconds) is OK
 - Client with lock can keep working
 - Locks are not lost

Reads in Paxos-based RSM

- Goals:
 - Linearizability: A read must see the effect of all accepted writes
 - Get read accepted to one of the slots in the replicated log

- Concurrent processing does not matter
 - Issue: Poor performance at scale
 - Alternative: serve reads at all replicas
 - Good Performance
- Approach:
 - For every key, every replica stores (value, version)
 - Return latest version out of majority of replicas
 - Issue: Different replicas may have different versions

Zookeeper

- General Info
 - Open-Source coordination service
 - Vs. Chubby
 - Similarities
 - Service and not library
 - Small file-system like hierarchy
 - Differences
 - Lower level API allows for more flexible coordination patterns
 - Weaker guarantees for better performance
 - Establishes core primitives rather than locks
 - Leader runs consensus protocol called Zab (Zookeeper atomic broadcast)
 - Use a weaker guarantee to avoid blocking as much as possible
 - Chubby cache invalidation leads to arbitrary delays due to waiting to hear back from all clients
 - *Almost* linearizability
 - Writes are strongly consistent and ordered
 - All requests from a single client are processed in FIFO order
- Implementing locks in Zookeeper
 - Clients register to watch a file
 - Zookeeper notifies client when file is updated
 - Zookeeper does not need to block for reads

- Zookeeper serves reads in any replica (vs only master in Chubby)
- Lock implementation
 - Acquire a lock by creating a file
 - If file already exists, watch for update
 - Upon watch notification, try to re-acquire the lock
 - Safe as writes are strongly consistent and ordered
 - Caused extremely bursty traffic as all processes all fight for lock
- Zookeeper writes do no invalidate caches
 - Zookeeper writes blocks for other writes
 - Writes hold the cost of computation
 - Blocking is necessary for strong consistency
- Optimizations
 - Client pipelines all updates in a FIFO so service can run them all at once

Lecture 11: CAP Theorem, Eventual Consistency, System Correctness, Conflict Free Replicated Datatypes

Strong consistency needs consensus (and blocking and order)

CAP Theorem: It is impossible for a distributed system to satisfy all 3 of the following:

- Consistency (Safety)
- Availability (Liveness)
- Partition Tolerance: (Unreliability)
 - CP and AP are most common
- Proof
 - Setup
 - Replicas in different network partitions cannot communicate
 - Partitions are transient and change
 - Responding without violating linearizability → Requires majority partition
 - Ensuring replicas always serve clients → Requires dropping consistency

PACELC

- If **P**artition: Choose **A**vailability vs **C**onsistency
 - Partitions are not the common case
- **E**lse: Choose **L**atency vs **C**onsistency
- Strong consistency means sacrificing latency

Order independent updates: Split the state of registers, clients can contact any replica

Eventual Consistency

- If no new updates, all replicas eventually converge to the same state
 - Liveness → Cannot be violated with a finite execution trace
 - Delivered updates might not be in the same state
- Issues
 - Lack of referential integrity
 - Updates may appear to be lost

- Thus, Amazon S3 uses read-after-write consistency after initial put. Most data is write-once

Strong Convergence

- Replicas that have delivered. The same set of updates have equivalent state
 - Differentiates between same set of updates vs delivered updates

Strong Eventual Consistency

- Heavily application dependent → Application must be able to tolerate disorder
- Limited on reads

Techniques for Checking System Correctness

- Testing: Given a system, add set of assertions and checks
 - Lightweight and automatic, but limited guarantees
- Monitoring/Runtime Verification: Create a monitor system, run system in production and check monitor output
 - Lightweight and automatic, but limited guarantees and added overhead to production
- Model Checking: Model the system as a state machine, come up with a set of properties
 - Guarantees and automatic, but expensive and white box
- Deductive Verification: Come up with a set of properties and construct a proof that system satisfies properties
 - Strong guarantees, but manual effort and white box

Facebook TAO

- Social Graph Database:
 - Nodes are people, groups, pages
 - Edges are follows, friend relations
- Linearizability
 - Use monitoring to gather 12 days of execution traces
 - Generate logs for each processed request
 - Check if execution is linearizable

- Issue: Validity due to Clock Synchronization
 - Comparison of time across different machines
 - Solution: Assume 99th percentile skew of 35ms
 - 0.003% violations in linearizability

Strong Convergence

- Replicas that have delivered the same set of updates have equivalent state
 - Issue: Does not care about the resulting state
 - Solution: Require an algorithm for conflict resolution

CDRTs (Conflict Free Replicated Datatypes)

- Guarantee strong eventual consistency (as long as same updates are applied)
 - Registers, counters, sets, lists
- Operation-based synchronization
 - Each replica broadcasts all update operations
 - Replicas apply updates from clients once they are received
 - Requires updates to be commutable
 - Less communication overhead
 - Requires reasoning over partial orders of operations
- State-based synchronization
 - After each update, each replica broadcasts its state S to other replicas
 - Replicas execute merge every time they receive another replica's state
 - Merge function is customized: $S \times S \rightarrow S$
 - Simpler to reason about
 - States exchanged may be large

Lecture 12: Sharding, Data Lookup, Chord

Assumptions thus far

- Every replica has to handle every operation
 - Repeated work means we can only scale up to the slowest replica
- Every replica has to store the complete state
 - We can only support state size up to storage of smallest replica
- Solutions:
 - Vertical scaling: Upgrade the worst machine
 - Horizontal Scaling: Get more machines, fix protocols to adapt to new machines
 - Cheaper

Sharding

- Idea: Partition state of one machine across multiple machines
 - Statistical hashing for less key skew
- Issue: Not fault tolerant, if shard fails, the state is lost forever
 - Solution: Create replicas and run Paxos for each shard
 - Does NOT improve scalability

Centralized Data Lookup

- Idea: Have a lookup service to know the location of each item
 - Hadoop File System uses a Name Server and Data Server
- Contacting the lookup service:
 - When a new server joins (server learns items it needs)
 - When we have a new item (service assigns it)
- Straightforward protocol, but hierarchical and thus not scalable

Decentralized Data Lookup

- Edge Cases
 - Randomly assign data to nodes
 - Even distribution but hard to find items
 - Place everything on one machine

- Easy to find items, but uneven distribution
- Naïve Approach: Assign key to nodes based on ranges
 - More even, but request distribution can still be uneven
 - Not all keys are used the same
- Hashing: Take a value and move it to an output space
 - Should minimize collisions
- Modulo Hashing: Compute modulo via number of servers
 - Very efficient but can lead to key skew
 - Additionally, changing machines requires moving all keys
 - $\text{Keys} / \text{Nodes}$ is minimum number of keys to move
 - Ideally, K is large and N is small

Consistent Hashing

- Idea: Represent hash output space as a circle
 - Every node ID is also hashed and points to a point in the circle
 - Note: Hash is not equi-distance/uniform
 - Thus, some servers take double the keys
- Keys are partitioned across servers
 - To locate a key, hash the key and find its successor server
- Adding/Removing Nodes
 - Move K/N keys
 - Minimizes migration of state upon change in set of servers
 - When removing a node, the successor takes over the shard

Virtual Nodes

- Each server gets multiple random IDs
 - More IDs mean better load balancing
 - Each ID corresponds to a virtual node
 - Vary id numbers across servers (bigger node gets more)
- When adding a node, you need to move data across nodes

Data Lookups with Shards

- Successor Pointers
 - Each node knows who its successors are
 - If you do not have value for key, forward to successor
 - $O(N)$ lookup
- Finger Table
 - Idea: Maintain $\log(N)$ pointers to reach $O(1)$ lookup
 - i 'th entry at node n points to successor of $\text{hash}(n) + 2^i$
 - number of entries = number of bits in hash value
 - Use binary lookup tree rooted at every node to find closest successor to value
 - Node Joins/Removals: All nodes know their predecessor
 - Update n 's finger table
 - Update all other nodes' finger tables
 - Move data around accordingly
 - Periodically maintain successors and predecessors using a stabilization protocol
 - Fault Tolerance
 - Replace crashed node data at its closest successors
 - Maintain a successor list of size $\log(N)$

Lecture 13: Amazon Dynamo, Data Store

Data Store

- NoSQL
 - Easier to get started and prototype
 - Good for small set of keys that each request needs to access
 - Lighter and thus more efficient. Flexible
 - Pushes more work onto the application
 - No schema
- SQL
 - Highly optimized for joins and other SQL constructs
 - Strong Consistency guarantees
 - Harder to implement

Service Driven Applications

- Frontend is ideally stateful
- Backend is ideally stateless (used to compute aggregates)

Service Level Agreements (SLAs)

- Contracts that describe service performance
 - Represents distribution of response times given load
- Fan out application:
 - End to End SLA Percentile:
 - For 3 backends, 90th percentile 90ms
 - $0.9 * 0.9 * 0.9 = 0.73$
 - End to End SLA = 73 percentile 100ms
 - Parallel requests to backend that get aggregated
 - More services entail harder guarantees for tail

Amazon Dynamo: Highly available distributed data store

Dynamo Motivation

- Datacenters across the world composed of thousands of servers and network components
 - Components fail continuously
 - Requires persistent state given failures

Dynamo Goals

- Sacrifice Consistency to maximize Availability

Dynamo Big Picture

- Provide Key/Value API (NoSQL) instead of SQL interface
 - No operations span across multiple keys
- Weak Consistency
 - Strong consistency not possible with partitions and availability
- Conflict resolution (for concurrent requests)
 - Resolved during reads → Reads uncover conflicts
 - Writes are never rejected, users can proceed if all is well
 - Dynamo supports Application and Data store conflict resolution
 - Application allows for hand-tuned conflict resolution at the cost of developer overhead
 - Data store is automatic but not natural

Dynamo Design and Implementation

Problem to Address	Solution
Partitioning	Consistent Hashing
High Availability for writes	Vector Clocks
Consistency	Sloppy Quorum
Temporary Failures	Hinted Handoff
Permanent Failures	Merkle Trees
Membership and Failure Detection	Gossip-based Membership

Dynamo Decentralization: All nodes do same work

- Helps with provisioning, maintenance, scaling, and reliability

Dynamo uses Virtual Nodes for Consistent Hashing

- Creates uniform data distribution, awareness to node heterogeneity

Dynamo Replication

- Replicate value for every key at N nodes in clockwise fashion
 - If node is a duplicate physical node, skip it
- Execution of writes
 - Write received by coordinator
 - Coordinator forwards to successors
 - Declare success if R reads or W writes succeed

Data Look Up

- Any server can receive and process requests
 - If server receives a key request, if it doesn't have the key, ask the next node ($O(N)$ worst time)

Gossip

- Once per second, each server contacts a randomly chosen other server
 - Both server exchange their lists of known servers
 - If servers do not change frequently, results in $O(1)$ Lookup time

Consistency via Sloppy Quorums

- Rather than consensus, select any available server to handle requests
 - N replicas for every key
 - Serve reads and writes by forwarding request to first N-1 reachable successors
 - Eventual Consistency but temporary inconsistencies
 - System supports custom merge operations

Vector Clocks

- Naïve Approach: Store a vector clock with each key-value pair
 - Not scalable, only a few nodes can write to item
- Dynamo Approach: Dynamic list of (coordinator node, counter) pairs
 - Determines if a version supersedes another

- Tracks causal relationships between updates
 - Can check for concurrent requests, performs reconciliation via a custom merge function
- Get(key) → List of <value, context> pairs
 - When responding to GET, Dynamo always returns the vector clock for the value returned
- Put(key, value, context)
 - Context indicates which versions this version supersedes or merges
- Clock Truncation
 - Case: More nodes processes PUTs to same key (after fault or network partition)
 - Dynamo stores the time of modification with each version vector entry
 - When version vector is more than 10 nodes long, Dynamo drops the least recently processed key
 - Causes false concurrency
- Comparisons
 - Vs Chord
 - Simpler protocol, worse “worst case” scenarios, but those never come up in practice
 - Chubby/Zookeeper
 - Both are services for users
 - Dynamo is fault tolerant data store, Chubby helps build one
 - Dynamo is not strongly consistent and surfaces inconsistencies to clients

Lecture 14: Failures in Dynamo, Merkle Trees, Two Phase Locking

Failures in Dynamo

- Temporary: Node is unreachable temporarily or due to network partition
 - Issue: Suppose data must be written to 3 nodes to be successful
 - If one node is unreachable, we must either delay the response or respond without durability
 - Solution: Hinted Handoff
 - Coordinator tries writing to subsequent nodes (beyond N successors of key)
 - Coordinator informs recipient of the intended node (which is currently down)
 - Recipient keeps “hint” in its metadata
 - Once intended recipient is back, we forward the key they should have and delete as they are adequately replicated
 - Issue: Hinted Handoff node may crash before it finishes replicating data
 - Solution: Replica synchronization
 - Nodes nearby on ring periodically gossip
 - Compare (k, v) pairs, copy missing keys
 - We need to compare all key pairs
- What if $W=1$?
 - High availability, no durability. If node fails, we lose access to some data
- If $W > 1$ and whole data center fails, we lose data. Thus, we construct preference list to span multiple data centers
- Permanent: Node is down permanently and won't return
 - Requires rebalancing by manual adding/removing nodes

Efficient Synchronization with Merkle Trees

- Merkle tree hierarchically summarizes key value pairs a node holds
 - Leaf holds hash of one key-value pair
 - Internal node = hash of concatenation of children hashes

- Comparing roots: if match, values match
 - If no match, compare children. Iterate down
- Finding differences (Reconciliation)
 - Exchange and compare hash nodes from root downwards, prune when hashes match

Local vs Global Failure Detection

- Most failures in Dynamo are temporary
 - Thus, Dynamo does not gossip failure information and stores it locally

Performance

- Tail is much higher than average
- Real workloads are diurnal
- Writes take more than reads (due to disk access)
- Some applications need even lower 99.9 writes
 - Use pipelining: Each node pushes to local in-memory buffer
 - Sacrifices durability (if node crashes, buffer is lost)

Consistent Hashing with Merkle Trees

- Copying new data over requires recomputing Merkle Trees
- Solution: Partition space into statically fixed number of equal sized shards
 - Place shard on first N virtual nodes after its end
 - Thus, we only need to identify which shards to hand off when adding node

Sharding Coordination

- Idea: We must enable load balancing across partitions
 - If operations touch multiple partitions

Atomicity of Transactions (Sequences of Operations)

- All-or-nothing → However, can cause some operations to be lost
- Goal:

- Serializability: For each execution of the system, there exists an equivalent execution that is fully serial
- Solution: Two Phase Locking
 - Apply locks on all necessary shards before attempting transaction
 - Algorithm
 - Client submits transaction to a coordinator (Transaction Coordinator)
 - TC acquires locks on all data involved
 - Once locks acquired, execute transaction and release locks
 - Ideas:
 - Disjoint transactions can execute concurrently
 - Related transactions wait on each other
- Alternative: Snapshot Isolation
 - Each transaction operation on a snapshot of the database can't read different data items from different times
 - Issue: Transaction cannot succeed if any partition becomes unavailable
 - Solution: Every partition should be implemented as an RSM → Higher availability

Lecture 16: Guest Lecture, Beaver, Cuttlefish, OrbWeaver

Partial Distributed Snapshots (Beaver – OSDI 2024)

- Challenge: Classic snapshot protocols (e.g., Chandy-Lamport) assume full control of all participants and closed-world causality — unrealistic in cloud microservices with external calls and partial observability.
- Solution:
 - Beaver introduces Partial Snapshots that capture causally consistent views even with unknown external interactions.
 - Uses Monolithic Gateway Marking (MGM) by repurposing existing SLBs (Software Load Balancers).
 - Optimistic Gateway Marking (OGM) handles coordination with multiple asynchronous SLBs.
 - Enables fast, zero-overhead, non-blocking snapshots suitable for debugging, analytics, and GC.

Fair Cloud Execution for Financial Exchanges (Cuttlefish – WIP)

- Challenge: Traditional exchanges achieve fairness via physical proximity and cable length. In the cloud, latencies are variable and paths are unknown.
- Solution:
 - Cuttlefish introduces a Virtual Time Overlay, where every operation is quantized in virtual time.
 - Ensures communication and computation synchrony through frozen/advanced virtual cycles, independent of real-time variability.
 - Provides fair and deterministic execution across heterogeneous systems and network conditions.

Opportunistic Failure Detection (OrbWeaver – NSDI 2022)

- Challenge: Traditional heartbeats trade off detection speed for overhead; message loss and failures are indistinguishable.
- Solution:

- OrbWeaver uses weaved streams, inserting tiny metadata packets into natural inter-packet gaps (on the order of 100ns).
- Guarantees high-fidelity, low-cost in-band telemetry and coordination (e.g., clock sync, failure detection) with minimal power/latency overhead.
- Uses programmable switches (e.g., Tofino) to achieve detection times $< 1\mu\text{s}$ at 100 Gbps.

Lecture 17: Cloud Native Applications, Microsoft Azure

Storage Types

- Blobs
 - Unstructured files with binary data
 - Media, log, text files
- Tables
 - Structured with schema – each row is a data item
 - User profiles, business data
- Queues
 - Message Delivery
 - Orchestrate Workflows

Cloud Native Applications

- Applications are composed of services
 - I/O data using blob
 - Intermediate and final analysis on tables
 - Orchestration and workflows with queues
- Idea: Cloud applications are orchestrated using workflows
 - Fan Out, Fan In, Timer
 - Developers don't deal with development
 - Providers can multiplex
 - Makes optimization and analysis hard

Windows Azure Storage

- Main Ideas
 - Global Namespace → finds data in huge storage
 - Uses DNS, Data is accessible through URI
 - All items in a transaction MUST be in the same partition
 - Disaster recovery: replication across data servers
 - Primary Backup Replication
 - Scalability and load balancing → Serve many users

- Multi tenancy → Dedicated storage leads to waste

Cloud Architecture

- Separation of services and clusters
 - Service/application deals with application tasks
 - Cluster managers handle other tasks

Azure Storage Architecture

- Storage Stamp: Cluster of storage nodes, one per account
 - Find account stamp using the location service
 - Stream Layer
 - Distributed file system, manages extents durably
 - Only allows reads and appends
 - Stream: Extent pointer seq
 - Cheap to construct, logical
 - Extend: Append block seq
 - Expensive to construct, physical
 - Sealed means no more appends
 - Append many small items together until they fill up an extent
 - Reversal of extents is fast due to modifying list of pointers
 - Stream Manager: Paxos replicated service, not on data path
 - Extend Node: Actual data servers, handles replication
 - Partition Layer
 - Maps abstractions into streams, manages transactions and consistency, handles scaling through sharding
 - Idea: Stateful operations are saved in a log via range partitions
 - Keeps track of materialized checkpoint of state and rest of commit log
 - Each partition is served by a different partition server

- Partition Manager: Assigns partitions to Partition Server, load balances
- Partition Server: Transactional guarantees for partition
- Lock Service: Paxos lock service, Election for PM, Lease for PS to serve partitions
- Range Partition
 - Storage Stream
 - Metadata stream: All information to load a partition to a server
 - Commit log stream: Holds all operations since last checkpoint
 - Row Data Stream: Stores checkpoints of data
 - In Memory State
 - Memory table: All updates since last checkpoint
 - Row data cache: In memory version of checkpoints
- Workflow
 - Partition server starts serving a range partition by loading it
 - Update operations go to commit log and memory table
 - Reads first go into memory log and then to in memory checkpoint
 - When memory table or commit log reach a limit, make a checkpoint
- Pros
 - Optimized for streams
 - No need to replay entire operation log from beginning
 - No full checkpoint needed to recover
- Load Balancing
 - Partition Server checkpoints range partition
 - Stops serving traffic
 - Partition Server creates new streams
 - Very fast as we just change pointers
 - Modify new partition metadata streams

- Start serving requests
- Frontends
 - Stateless routers that handle authentication and data access
 - Optimizations:
 - Cache frequently accessed data
 - Streams data directly from the stream layer when possible
- Goal: Use stamps at 70%
- Stamp Replication
 - Intra Stamp
 - Synchronous, on critical path of request, durable
 - Low latency
 - Inter Stamp
 - Off critical path on per object basis, for disaster recovery
 - Optimal use of network bandwidth
- Eraser Coding
 - Reduces storage cost post-extent sealing, increases read latency
 - For any extent, store one original copy
 - Store (M+N) erasure coded chunks

Lecture 18: Spanner

Goals

- Scalable
- Multiversion
- Globally Distributed → Reliable
- Synchronously Replicated → Fault Tolerance and Locality
- Distributed/Cross Partition Transactions

Guarantees

- Externally consistent transactions (Strict serializability)
 - Also called real-time ordering guarantee
 - Bans reordering of non-overlapping transactions
- Snapshot isolation transactions
 - No need to lock, just serve reads for specific timestamp

Cross Shard Geodistributed Transactions

- Idea: Commit timestamps for transactions using TrueTime API
- TrueTime API
 - API exposes real-world clock uncertainty across machines
 - Guarantees depend on uncertainty bounds
 - If large uncertainty, Spanner slows down and waits
 - Uncertainty is reported conservatively for correctness

Spanner Architecture

- Big Picture
 - Split up into zones
 - Zones are added, removed, and replicated
 - Datatype is (key, timestamp) → value
 - Many timestamps of same key live in the Database (multi version)
 - Data is split into tablets, which compose spanservers
 - Storage of files and commit logs

- Replication Management
 - Paxos RSM for each tablet
 - Long lived leader and time-based leases
 - Writes go to leader to start proposal
 - Reads are served from any up-to-date replica
- Concurrency Control
 - Each leader has a lock table
 - Transactional operations acquire locks
 - Other operations bypass the lock table
- Cross-zone Transactions
 - For local transactions, a lock table is enough
 - For cross-tablet, we need a two-phase commit
 - One leader of tablets is chosen as coordinator
 - Each leader has a transaction manager data structure
 - State of transaction manager is started with Paxos
 - Strong consistency for leader faults
- TrueTime API
 - Each machine knows an interval for the real time it can have
 - `TT.now()`, `TT.after(t)`, `TT.before(t)`
 - Idea: We can guarantee that $T1.start < T2.end$
 - Thus, we can establish causality
 - Implementation
 - GPS based timemasters, some atomic
 - Servers sync with timemasters
- Implementing Spanner guarantees
 - Paxos clusters
 - Replicas use Paxos for lease votes
 - Lease interval starts when leader discovers it has quorum, ends when some lease votes have expired
 - Lease votes are implicitly extended on successful writes to state
 - Correctness depends on lease disjointness

- Each Paxos leader interval is disjoint from past and future leaders (they cannot overlap)
 - Ensuring lease disjointness
 - Issue a Paxos write to write the lease interval
 - Next leader waits for all intervals to pass before starting its lease
 - Optimizations:
 - Each replica follows a single-vote rule: If it grants a lease at time T , it waits until $TT.after(T.end + lease_len)$ before it grants another
 - At least one replica will be common between any two consecutive leaders (Use quorum for lease votes)
- Read-write transactions
 - Spanner wants to provide strict serializability
 - Monotonically increasing commit timestamps per tx
 - Establishes a sequential order
 - Exploit disjointness to establish a monotonic timestamp across leaders
 - Real Time Ordering
 - If start of T_2 occurs after the commit of T_1 , then $T_2 > T_1$
 - Commit timestamp s is after coordinator's latest time
 - Wait until my earliest time has gone beyond s before writing
 - Implementation
 - Writes are buffered at the client until commit
 - Client issues reads to leaders of groups
 - To avoid deadlock, the locks have a timeout
 - When all reads are done, start a two phase commit
 - Two Phase Commit
 - First, client sends messages to all leaders
 - Non-coordinator (upon receiving)
 - Acquires write locks
 - Prepares timestamp T that is larger than any previously assigned

- Logs the prepare(T) record using Paxos
 - Sends prepare message to coordinator
- Coordinator leader (on receiving commit)
 - Acquires write locks
 - Chooses commit timestamp T_s that is greater than:
 - All prepare(T) timestamps
 - $TT.now().end$ at the time of receiving commit_tx
 - All previous assigned transaction timestamps
 - Logs the commit(T_s) record using Paxos
 - Wait until $TT.after(T_s)$
 - Send message to client and other leaders
 - Other leaders log the write and then release the locks
- Serving snapshot reads
 - No need for blocking
 - Each replica has a t_safe time and can only serve reads up to it
 - $T_safe = \min(t_paxos, t_tm)$
 - T_paxos = timestamp of highest applied Paxos write
 - T_tm depends on prepared transactions

Lecture 19: Microservices, Mucache

Microservices

- General Info
 - Splits functionality into multiple services
 - Issue: They can be written in multiple languages with different data representations
 - Solution: Serialization Libraries
 - Offers unified representation across frontend implementations of data
 - Communication between services using established APIs and RPC
 - Each service might have its own state and are deployed in their own machines
- Microservice Application Requirements
 - Need to find which other services exist in the network and how to call them
 - Authenticate requests from other services
 - Serialization/Deserialization
 - Encrypt requests for privacy
 - Monitor request-level metrics
 - Often implemented using a Service mesh (application independent)
- Service Mesh
 - Used to build microservices by being its functionality
 - Come together via a proxy/sidecar
 - Sidecar intercepts all incoming and outgoing communication
- Microservice Benefits
 - Mostly organizational
 - Allows for easier delegation of work into teams
 - If service fails, we can simply call the team versus the entire development staff
- Microservice Graphs
 - Note: edges imply one or more requests
 - Also, each service can be deployed on many nodes
 - Sharded for scalability

- Replicated for fault tolerance
- Shortcomings
 - Performance worsens from sidecars
 - Data goes twice through network per request
 - Programmability, Correctness
- Optimizations
 - Improve performance of stack (serialization algorithms)
 - Running sidecar on the same VM
 - Colocating services
 - Caching

Caching

- Service calls without side effects can be cached and not remade
 - If downstream data has not been modified
- Issue: Coherence: Database updates affect validity of upstream caches
- Solutions
 - Backend only Caching
 - Custom invalidation protocols per application
 - Give expiry to cache items(Time-to-Live eviction per service)
 - Ambiguous to set the time of TTL

MuCache

- Automatic cache invalidation for microservice applications
- Goals
 - Correctness: Behaves the same as original
 - Determined using traces
 - Create traces of client requests. Should be in the original set of traces
 - Sometimes requires linearizability as subset is an impossible execution (needs synchronization)
 - Clients do not need real time ordering

- Given a set of cache-enabled traces, there exists a reordering such that the new order is in the original
 - Low Overhead: Cache management off the critical path
 - Support Real Microservice Deployments: Sharding and dynamic call graphs
 - Plug and Play
 - Incremental Deployment: Support evolving graphs where MuCache is not deployed in all services
- Architecture
 - Developer declares read only methods
 - MuCache Components
 - Cache
 - Wrappers (Intercepts communication to database)
 - Library that wraps all calls to services and database
 - Adds messages to cache manager queue
 - Startcall, Endcall, Invalidatecall
 - Keeps total order of call starts, call ends, invalidations
 - If a call's results are invalidated by a later write, then the cache item upstream is evicted
 - Cache Manager (invalidates caches)
 - One thread per shard
 - Has Input message (FIFO) queue
- MuCache Invalidation
 - Diamond patterns
 - S1 may become stale if S4 is updated without notifying S2 or S3
 - Issue: Invalidation might overtake the save and create divergence
 - Solution: Calls are invalidated unnecessarily
 - Cache manager cannot differentiate between write and read
 - Thus, we invalidate the write and assume all reads are batched before any writes
- Dependency Tracking
 - Each call keeps track of its visited services

- Before checking the cache at another request going through, we see that visited services intersects with cache results
 - Thus we do not use the cache